


Retrieval ranked list



Epistemic boundary: the model cannot reason over evidence it never sees.

"Not in top_k" $\neq$ "not relevant" $\neq$ "does not exist" — OWA (Ch4).

Conversely: too large a top_k $\rightarrow$ drowned in noise/distractors $\rightarrow$ the LLM gets "lost" in context, quality drops rather than rises. top_k is an intent-dependent policy, not a constant.